
Accelerating the Convergence of Policy Gradient Methods

Anant Shyam **Daniel Cao**
Cornell University Cornell University
ais64@cornell.edu dyc33@cornell.edu

Abstract

1 Policy gradient methods are critical for finding the optimal parameters for policies
2 in order to be effective in reinforcement learning (RL) settings. However, such
3 methods are often quite expensive and incur large variance estimates of the gradient
4 of the performance function. We provide a detailed analysis of various approaches
5 to accelerating the convergence of policy gradient to an ϵ -stationary point of the
6 optimal policy, as well as giving a high level discussion of the proof techniques for
7 deriving the bounds relating sample complexity and convergence.

8 1 Introduction

9 In this paper, we'll analyze the following previous works: Sample Efficient Policy Gradient Methods
10 with Recursive Variance Reduction (Pan Xu, Felicia Gao & Quanquan Gu), Momentum-Based Policy
11 Gradient Methods (Feihu Huang, Shangqian Gao, Jian Pei & Heng Huang), and Policy Gradient
12 Method For Robust Reinforcement Learning (Yue Wang, Shaofeng Zou). We'll refer to these papers
13 as SRVR-PG, MBPG, and RPG respectively.

14 1.1 RL Background

15 Through this project, we'll be mainly considering infinite horizon MDPs, i.e $\mathcal{M} = \{S, A, \mathcal{P}, r, \gamma, \rho\}$,
16 where S is the state space, A is the action space, \mathcal{P} is the transition function such that $\mathcal{P}(s'|s, a)$
17 is the probability of transitioning to a state s' given the previous state s and action a . r represents
18 a reward function, such that $r(s, a) \in [-R, R]$, where R is a constant. γ represents the discount
19 factor, and $\gamma \in (0, 1)$, and ρ represents the distribution over the initial state s_0 . Define a policy π_θ ,
20 parametrized by the parameter θ , such that $\pi_\theta : S \rightarrow A$. Moreover, τ represents a trajectory.

21 We want to find a policy that maximizes the total expected cumulative reward, given by $R(\tau) =$
22 $\sum_{h=0}^{H-1} \gamma^h r(s_h, a_h)$, where H represents the horizon length, and $\tau = \{s_0, a_0, \dots, s_{H-1}, a_{H-1}\}$. Let
23 $J(\theta) = \mathbb{E}_\tau[R(\tau)]$. We define an stationary point of $J(\theta)$ to be θ such that $\|\nabla J(\theta)\|_2 = 0$ and
24 we will also define an ϵ -stationary point of $J(\theta)$ to be θ such that $\mathbb{E}[\|\nabla J(\theta)\|_2^2] \leq \epsilon$. Moreover,
25 assuming the Markov property holds, we'll define $p(\tau|\theta)$, the probability of observing a given
26 trajectory τ given the policy parameters θ to be equal to $\rho(s_0) \prod_{h=0}^{H-1} \pi_\theta(a_h|s_h) \mathcal{P}(s_{h+1}|s_h, a_h)$.

28 1.2 Policy Gradient Background

29 Given a performance function $J(\theta)$, it is often not possible to calculate $\nabla J(\theta)$ exactly. Since
30 $J(\theta) = \mathbb{E}_\tau[R(\tau)]$, we can say that $\nabla J(\theta)$

31
$$= \nabla_\theta \mathbb{E}_\tau[R(\tau)] = \nabla_\theta \int_\tau R(\tau) p(\tau|\theta) d\tau = \int_\tau R(\tau) \nabla_\theta p(\tau|\theta) d\tau$$

32
$$= \int_\tau R(\tau) (\nabla_\theta p(\tau|\theta) / p(\tau|\theta)) p(\tau|\theta) d\tau = \mathbb{E}_{\tau \sim p}[\nabla_\theta \log p(\tau|\theta) R(\tau)]$$

33
$$= \mathbb{E}_{\tau \sim p}[\nabla_\theta \log p(\tau|\theta) R(\tau)]$$

35 Note that we don't know the distribution p so we can't directly calculate this expectation.
 36 However, we can estimate the gradient through using a batch of trajectories τ_1, \dots, τ_N , where
 37 our estimated gradient $\hat{\nabla}_\theta J(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_\theta \log p(\tau_i|\theta) R(\tau_i)$. Since $\nabla_\theta \log p(\tau_i|\theta) = \nabla_\theta$
 38 $\log(\rho(s_0) \prod_{h=0}^{H-1} \pi_\theta(a_h|s_h) \mathcal{P}(s_{h+1}|s_h, a_h)) = \nabla_\theta \sum_{h=0}^{H-1} \log \pi_\theta(a_h^i|s_h^i) = \sum_{h=0}^{H-1} \nabla_\theta \log$
 39 $\pi_\theta(a_h^i|s_h^i)$, and by using the definition of $R(\tau_i)$, we can rewrite our gradient estimate to be
 40 $\frac{1}{N} \sum_{i=1}^N (\sum_{h=0}^{H-1} \nabla_\theta \log \pi_\theta(a_h^i|s_h^i)) (\sum_{h=0}^{H-1} \gamma^h r(s_h^i, a_h^i))$. We will denote $g(\tau_i|\theta) = (\sum_{h=0}^{H-1} \nabla_\theta$
 41 $\log \pi_\theta(a_h^i|s_h^i)) (\sum_{h=0}^{H-1} \gamma^h r(s_h^i, a_h^i))$.
 42

43 2 Robust Policy Gradient Analysis

44 2.1 Algorithm Overview

45 In this paper, the first robust policy gradient method was developed, which was proven to convergence
 46 to an ϵ optimal policy in $\mathcal{O}(\epsilon^{-3})$ complexity. In this paper, we consider the following MDP:
 47 (S, A, P, c, γ) , where S, A are defined as usual, P is a transition kernel, where p_s^a is the distribution
 48 over the next states given s, a , γ is the discount factor, and c is the cost function, $0 \leq c(s, a) \leq 1$.
 49

50 Define a sequence of transition kernels $\kappa = (P_0, P_1, \dots)$. The robust value function is de-
 51 fined as $V^\pi(s) = \max_{\kappa} \mathbb{E}_\kappa[\sum_{t=0}^{\infty} \gamma^t c(s_t, a_t) | s_0 = s, \pi]$. Since we are trying to minimize the
 52 cumulative discounted cost, the robust value function is the worst case discounted cost with policy
 53 π and starting state s . Now, let the performance function $J_\rho(\pi) = \mathbb{E}_{s \sim \rho}[V^\pi(s)]$. In robust policy
 54 gradient, we aim to find θ such that the performance function is minimized. In other words, we are
 55 trying to solve the following optimization problem: $\operatorname{argmin}_{\theta \in \Theta} J_\rho(\theta)$.
 56

57 Assume that the policy class Π_θ is differentiable and k_π Lipschitz continuous. If this assumption
 58 holds, then we can say that $\psi_\rho(\theta) = \frac{\gamma R}{(1-\gamma)(1-\gamma+\gamma R)} \sum_{s \in S} d_{s_\theta}^{\pi_\theta}(s) \sum_{a \in A} \nabla \pi_\theta(a|s) Q^{\pi_\theta}(s, a) +$
 59 $\frac{1}{1-\gamma+\gamma R} \sum_{s \in S} d_\rho^{\pi_\theta}(s) \sum_{a \in A} \nabla \pi_\theta(a|s) Q^{\pi_\theta}(s, a)$. Shown below is the pseudocode for robust policy
 60 gradient:
 61

```

62   for  $t = 0, \dots, T-1$  do
63        $\theta_{t+1} = \Pi_\Theta(\theta_t - \alpha_t \psi_\mu(\theta_t))$ 
64   return  $\theta_T$ 
65
```

66 Note that for all most all $\theta \in \Theta$, $J_\rho(\theta)$ is differentiable, which implies that $\psi_\rho(\theta) = \nabla J_\rho(\theta)$.
 67 Therefore, for all differentiable θ , we are performing a gradient descent update, and projecting the
 68 result on to the parameter space Θ .
 69

70 Now, let's consider a version of the previous algorithm, known as smoothed robust policy gradient.
 71 We optimize over a smoothed, differentiable approximation of J_ρ denoted by $J_{\sigma, \rho}$ for some $\sigma > 0$.
 72 Instead of the maximum, we'll use the LogSumExp (LSE) operator instead, where $\text{LSE}(\sigma, V) = \frac{1}{\sigma}$
 73 $\log(\sum_{i=1}^d e^{\sigma V(i)})$, where $V \in \mathbb{R}^d$. Define $Q_\sigma^\pi(s, a) = c(s, a) + \gamma(1-R) \sum_{s' \in S} p_{s, s'}^a V_\sigma^\pi(s') + \gamma R$
 74 $\text{LSE}(\sigma, V_\sigma^\pi)$, and define $V_\sigma^\pi(s)$ analogously. Let $J_\sigma(\theta) = \sum_{s \in S} \rho(s) V_\sigma^\pi(s)$. Shown below is the
 75 pseudocode for smoothed robust policy gradient:
 76

```

77   for  $t = 0, \dots, T-1$  do
78        $\theta_{t+1} = \Pi_\Theta(\theta_t - \alpha_t \nabla J_\sigma(\theta_t))$ 
79   return  $\theta_T$ 

```

80 . Note that we can compute that $\nabla J_\sigma(\theta) = B(\rho, \theta) + \frac{\gamma R \sum_{s \in S} e^{\sigma V_\sigma^{\pi_\theta}(s)} B(s, \theta)}{(1-\gamma) \sum_{s \in S} e^{\sigma V_\sigma^{\pi_\theta}(s)}}$, where $B(s, \theta) =$
 81 $\frac{1}{1-\gamma+\gamma R} \sum_{s' \in S} d_{s'}^\pi(s') \sum_{a \in A} \nabla \pi_\theta(a|s') Q_\sigma^{\pi_\theta}(s', a)$. The authors then claim the following theorem
 82 regarding this smoothed objective J_σ : $J_\sigma(\theta) - J_\sigma^* \leq C_{PL} \max_{\hat{\pi} \in (\Delta(|A|))^{|S|}} \langle \pi_\theta - \hat{\pi}, \nabla J_\sigma(\theta) \rangle +$
 83 $(\frac{\gamma R}{1-\gamma}) \frac{2 \log |S|}{\sigma}$, for some constant C_{PL} . This theorem really illustrates that using the smoothed
 84 objective, if the gradient $\nabla J_\sigma(\theta)$ is near zero (or small), then $J_\sigma(\theta)$ is within $\mathcal{O}(\frac{1}{\sigma})$ of the optimal

value J_σ^* . Choosing larger values of σ in this case and make this upper bound go to zero, showing optimality.

2.2 Convergence Analysis

Assume that the policy class Π_Θ is l_π smooth, and assume that $\|\nabla J_\sigma(\theta_1) - \nabla J_\sigma(\theta_2)\| \leq L_\sigma \|\theta_1 - \theta_2\|$. Now, we have the following convergence theorem:

For any $\epsilon > 0$, set $\sigma = \frac{2\gamma R \cdot \log|S|}{\epsilon(1-\gamma)}$ and $T = \frac{64|S|C_{PL}^2 L_\sigma C_\sigma}{\epsilon^2}$, then $\min_{t \leq T-1} J(\theta_t) - J^* \leq 3\epsilon$. This means that convergence to an ϵ optimum happens in $\mathcal{O}(\epsilon^{-3})$ steps. Now, let's prove this result. We'll summarize the results of this proof, due to this proof being quite involved.

We'll denote the gradient mapping as $G^\alpha(\theta) = \frac{1}{\alpha}(\theta - \Pi_{(\Delta A)^{|S|}}(\theta - \alpha J_\sigma(\theta)))$, and set all $\alpha_t = \frac{1}{L_\sigma}$. Now, we can note that $\min_{0 \leq t \leq T-1} \|G^\alpha(\theta_t)\| \leq \sqrt{\frac{2L_\sigma(J_\sigma(\theta_0) - J_\sigma^*)}{T}}$.

Now, we can bound by $\min_{t \leq T-1} J_\sigma(\theta_t) - J_\sigma^* \leq C_{PL} \max_{\hat{\pi} \in (\Delta(|A|))^{|S|}} \langle \pi_{\theta_W} - \hat{\pi}, \nabla J_\sigma(\theta_W) \rangle + (\frac{\gamma R}{1-\gamma})(\frac{2\log|S|}{\sigma})$, where $W = 1 + \arg\min_{t \leq T-1} \|G^\alpha(\theta_{W-1})\|$.

Through bounding $\max_{\hat{\pi} \in (\Delta(|A|))^{|S|}} \langle \pi_{\theta_W} - \hat{\pi}, \nabla J_\sigma(\theta_W) \rangle \leq \sup_{\hat{\pi} \in (\Delta(|A|))^{|S|}} \|\hat{\pi} - \pi_{\theta_W}\| \frac{\epsilon}{2\sqrt{|S|}C_{PL}}$ and since $\|\pi_{\theta_1} - \pi_{\theta_2}\| \leq 2\sqrt{|S|}$, our bound becomes $\frac{\epsilon}{C_{PL}}$. This implies that $\min_{t \leq T-1} J_\sigma(\theta_t) - J_\sigma^* \leq \epsilon + (\frac{\gamma R}{1-\gamma})(\frac{2\log|S|}{\sigma})$. Setting $\sigma = \frac{2\log|S|(\frac{\gamma R}{1-\gamma})}{\epsilon}$ and $T = \frac{64|S|C_{PL}^2 L_\sigma C_\sigma}{\epsilon^2}$ gives the upper bound of 3ϵ .

3 SRVR-PG Analysis

3.1 SRVR-PG Algorithm Discussion

Using a batch of trajectories to approximate the gradient $\nabla J(\theta)$ can still result in a large variance. To address this limitation, the SRVR-PG algorithm was developed. This algorithm consists of S epochs, denoted by $s = 0, \dots, S-1$. Denote the parameter of a reference policy to be $\tilde{\theta}^0 = \theta_0$ at the initialization of this algorithm. For each iteration s , N trajectories τ_1, \dots, τ_N are sampled from the reference policy $\pi_{\tilde{\theta}^s}$. Then, using these trajectories, we'll compute a gradient estimator $v_0^s = \frac{1}{N} \sum_{i=1}^N g(\tau_i | \tilde{\theta}^s)$.

We will now define a function $\mathcal{P}_\Theta(\theta)$, which we'll use to update the reference policy parameter θ^s . The update will be of the following form: $\theta_{t+1}^{s+1} = \mathcal{P}_\Theta(\theta_t^{s+1} + \eta v_t^{s+1})$, where v_t^{s+1} is our gradient direction. We let $\mathcal{P}_\Theta(\theta) = \arg\min_{u \in \mathbb{R}^d} [\mathbb{1}_\Theta(u) + \frac{1}{2\eta} \|\theta - u\|_2^2]$, where $\mathbb{1}_\Theta(u)$ is an indicator function representing whether u is in the parameter set Θ . We do the update $\theta_{t+1}^{s+1} = \mathcal{P}_\Theta(\theta_t^{s+1} + \eta v_t^{s+1})$ for $m-1$ epochs. Let t represent these epochs, $t = 1, \dots, m-1$. In each epoch we sample B trajectories τ_1, \dots, τ_B , and do the following gradient update: $v_t^{s+1} = v_{t-1}^{s+1} + \frac{1}{B} \sum_{j=1}^B (g(\tau_j | \theta_t^{s+1}) - g_\omega(\tau_j | \theta_{t-1}^{s+1}))$, and then do the update with \mathcal{P} to get the next parameter θ_{t+1}^{s+1} . After we run the above procedure for all $m-1$ epochs, to get an iterate θ_m^{s+1} . We then set our reference policy parameter $\tilde{\theta}^{s+1} = \theta_m^{s+1}$. Shown below is the pseudocode for the algorithm:

```

 $\tilde{\theta}^0 = \theta_0 \in \Theta$ 
for  $s = 0, \dots, S-1$  do
     $\theta_0^{s+1} = \tilde{\theta}^s$ 
    Sample trajectories  $\tau_1, \dots, \tau_N$  from  $p(\cdot | \tilde{\theta}^s)$ 
     $v_0^{s+1} = \hat{\nabla} J(\hat{\theta}^s)$ 
     $\theta_1^{s+1} = \mathcal{P}_\Theta(\theta_0^{s+1} + \eta v_0^{s+1})$ 

```

132 **for** $t = 1, \dots, m - 1$ **do**
 133 Sample B trajectories τ_1, \dots, τ_B from $p(\cdot|\theta_t^{s+1})$
 134 $v_t^{s+1} = v_{t-1}^{s+1} + \frac{1}{B} \sum_{j=1}^B (g(\tau_j|\theta_t^{s+1}) - g_\omega(\tau_j|\theta_{t-1}^{s+1}))$
 135 $\theta_{t+1}^{s+1} = \mathcal{P}_\Theta(\theta_t^{s+1} + \eta v_t^{s+1})$
 136 $\tilde{\theta}_{t+1}^{s+1} = \theta_{t+1}^{s+1}$
 137 Return randomly chosen $\theta \in \{\theta_t^s\}$ for all t, s

138 3.2 SRVR-PG Convergence Analysis

139 Assume that $\|\nabla_\theta \log \pi_\theta(a|s)\| \leq G$, and $\|\nabla_\theta^2 \log \pi_\theta(a|s)\| \leq M$ for all s, a . Assume that
 140 $\text{Var}(g(\tau|\theta)) \leq \xi^2$ for some constant $\xi > 0$. Moreover, assume that $\text{Var}(\omega(\tau|\theta_1, \theta_2)) \leq W$,
 141 where W is finite, and $\omega(\cdot|\theta_1, \theta_2) = p(\cdot|\theta_1)/p(\cdot|\theta_2)$. Denote the gradient mapping
 142 $\mathcal{G}_n(\theta) = \frac{1}{\eta}(\mathcal{P}_\Theta(\theta + \eta \nabla J(\theta)) - \theta)$, which conceptually represents the generalized projected gradient
 143 with respect to θ .

144
 145 **Convergence Theorem:** If the assumptions above hold, and if $\eta \leq \frac{1}{4L}$, and epoch size m
 146 plus mini-batch size B such that $B \geq \frac{72mnG^2(2G^2/(M+1)(W+1)\gamma)}{(1-\gamma)^2}$. Then, we get the following result:

147 $\mathbb{E}[\|\mathcal{G}_n(\theta)\|_2^2] \leq \frac{8(J(\theta^*) - J(\theta_0) - \mathbb{1}_\Theta(\theta^*) + \mathbb{1}_\Theta(\theta_0))}{\eta Sm} + \frac{6\xi^2}{N}$, where $\theta^* = \text{argmax}_{\theta \in \Theta} J(\theta)$. As proving
 148 this claim is extremely involved, we will give a high level overview of the all the proofs in the
 149 report.

150
 151 First, by the definition of \mathcal{P}_Θ , we have that $\theta_{t+1}^{s+1} = \text{argmin}_{u \in \mathbb{R}^d} \mathbb{1}_\Theta(u) + \frac{1}{2\eta} \|u - \theta_t^{s+1}\|_2^2 + \langle v_t^{s+1}, u \rangle$.
 152 We can also define the respective gradient mapping to be $\tilde{\mathcal{G}}_t^{s+1} = \frac{1}{\eta}(\theta_{t+1}^{s+1} - \theta_t^{s+1}) =$
 153 $\frac{1}{\eta}(\mathcal{P}_\Theta(\theta_t^{s+1} - \eta v_t^{s+1}) - \theta_t^{s+1})$. Using the convexity of $\mathbb{1}_\Theta(\cdot)$, $J(\theta)$ being L -smooth, and
 154 $\Phi(\theta) = J(\theta) - \mathbb{1}_\Theta(\theta)$, we can conclude that $\Phi(\theta_{t+1}^{s+1}) \geq \Phi(\theta_t^{s+1}) - \frac{\eta}{2} \|\nabla J(\theta_t^{s+1}) - v_t^{s+1}\|_2^2 +$
 155 $\frac{\eta}{8} \|\mathcal{G}_n(\theta_t^{s+1})\|_2^2 - \frac{\eta}{4} \|\mathcal{G}_n(\theta_t^{s+1}) - \tilde{\mathcal{G}}_t^{s+1}\|_2^2 + (\frac{1}{4\eta} - \frac{L}{2}) \|\theta_{t+1}^{s+1} - \theta_t^{s+1}\|_2^2$.

156
 157 Doing some analysis on $\bar{\theta}_{t+1}^{s+1} = \text{prox}_{\eta \mathbb{1}_\Theta}(\theta_t^{s+1} + \eta \nabla J(\theta_t^{s+1}))$ yields that $\Phi(\theta_{t+1}^{s+1}) \geq$
 158 $\Phi(\theta_t^{s+1}) - \frac{3\eta}{4} \|\nabla J(\theta_t^{s+1}) - v_t^{s+1}\|_2^2 + \frac{\eta}{8} \|\mathcal{G}_n(\theta_t^{s+1})\|_2^2 + (\frac{1}{4\eta} - \frac{L}{2}) \|\theta_{t+1}^{s+1} - \theta_t^{s+1}\|_2^2$. Now,
 159 when further expanding out $\|\nabla J(\theta_t^{s+1}) - v_t^{s+1}\|_2^2$ and take expectations, we get that
 160 $\mathbb{E}[\|\nabla J(\theta_t^{s+1}) - v_t^{s+1}\|_2^2] \leq \frac{1}{B^2} \sum_{j \in \mathcal{B}_t} \mathbb{E}[\|g_\omega(\tau_j|\theta_{t-1}^{s+1}) - g(\tau_j|\theta_t^{s+1})\|_2^2] + \|\nabla J(\theta_{t-1}^{s+1}) - v_{t-1}^{s+1}\|_2^2$,
 161 where \mathcal{B}_t denotes the minibatch, and B represents its' size.

162
 163 Through some more analysis, we can show that $\mathbb{E}[\|\nabla J(\theta_t^{s+1}) - v_t^{s+1}\|_2^2] \leq$
 164 $\frac{C_\gamma}{B} \sum_{l=1}^t \|\theta_l^{s+1} - \theta_{l-1}^{s+1}\|_2^2 + \|\nabla J(\theta_0^{s+1}) - v_0^{s+1}\|_2^2$, for a constant C_γ . Now, using the
 165 variance assumption and some of the previous results, we can say that $\mathbb{E}_{N,B}[\Phi(\theta_m^{s+1})] \geq$
 166 $\mathbb{E}_{N,B}[\Phi(\theta_0^{s+1})] + \frac{\eta}{8} \sum_{t=0}^{m-1} \mathbb{E}_N[\|\mathcal{G}_n(\theta_t^{s+1})\|_2^2] - \frac{3m\eta\xi^2}{4N} + (\frac{1}{4\eta} - \frac{L}{2} - \frac{3m\eta C_\gamma}{2B}) \sum_{t=0}^{m-1} \|\theta_{t+1}^{s+1} - \theta_t^{s+1}\|_2^2$.
 167 Setting η and B appropriately, we get the following result:

168
 169 $\mathbb{E}[\|\mathcal{G}_n(\theta)\|_2^2] \leq \frac{6\xi^2}{N} + \frac{8(\Phi(\theta^*) - \Phi(\theta_0))}{\eta Sm}$, which concludes the proof.

171 3.3 Extension to Parameter Based Exploration

172 This paper also discusses how to extend SRVR-PG to a policy gradient algorithm using parameter-
 173 based exploration (PGPE) as well. It assumes that the parameter θ follows a set prior distribution ρ .

174
 175 Given this parameter ρ , we define the performance function to be $J(\rho) =$
 176 $\int_\theta \int_\tau p(\theta|\rho) p(\tau|\theta) R(\tau) d\tau d\theta$. The goal of PGPE is to find the prior distribution ρ^* such that $\rho^* =$
 177 $\text{argmax}_\rho J(\rho)$. A linear deterministic policy π is chosen of the following form: $\pi_\theta(a|s) = \delta(a - \theta^T s)$,
 178 and the distribution update rule at every step is given by $\rho_{t+1} = \rho_t + \eta \nabla_{\rho_t} J(\rho_t)$. The gradient
 179 estimates are similar to those of SRVR-PG. Shown below is the pseudocode of this algorithm.

180
 181 $\tilde{\theta}^0 = \theta_0 \in \Theta$

```

182 for  $s = 0, \dots, S - 1$  do
183    $\rho_0^{s+1} = \rho^s$ 
184   Sample  $N$  policy parameters  $\theta_1, \dots, \theta_N$  from  $p(\cdot | \rho^s)$ 
185   Sample one trajectory  $\tau_i$  from each policy  $\pi_{\theta_i}$ 
186    $v_0^{s+1} = \hat{\nabla} J_\rho(\rho^s)$ 
187    $\rho_1^{s+1} = \rho_0^{s+1} + \eta v_0^{s+1}$ 
188   for  $t = 1, \dots, m - 1$  do
189     Sample  $B$  trajectories  $\tau_1, \dots, \tau_B$  from  $p(\cdot | \theta_t^{s+1})$ , sample  $\tau_j$  from each  $\pi_{\theta_j}$ 
190      $v_t^{s+1} = v_{t-1}^{s+1} + \frac{1}{B} \sum_{j=1}^B (g(\tau_j | \rho_t^{s+1}) - g_\omega(\tau_j | \rho_{t-1}^{s+1}))$ 
191      $\rho_{t+1}^{s+1} = \rho_t^{s+1} + \eta v_t^{s+1}$ 
192   Return randomly chosen  $\rho \in \{\rho_t^s\}$  for all  $t, s$ 
193

```

194 3.4 Results

195 The SRVR-PG algorithm was evaluated on many classic reinforcement learning environments, and
 196 the results of these experiments are displayed in Figure 1. From these figures, we can see that
 197 SRVR-PG overall has higher returns than the other estimators, including SVRPG and GPOMDP.

198 4 Momentum-Based Policy Gradient Methods

199 4.1 Brief Background and Motivation

200 In general, RL optimizes a policy $\pi_\theta(a|s)$, parameterized by $\theta \in \mathbb{R}^d$, to maximize the expected
 201 cumulative reward $J(\theta)$: $J(\theta) = \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau)] = \int R(\tau) p(\tau|\theta) d\tau$, where $\tau = (s_0, a_0, s_1, \dots)$
 202 represents a trajectory, and $p(\tau|\theta)$ is its distribution. While policy gradient methods like REINFORCE
 203 provide an effective framework by estimating $\nabla J(\theta)$ using stochastic approximations: $\nabla J(\theta) =$
 204 $\mathbb{E}_{\tau \sim p(\tau|\theta)} [\nabla \log p(\tau|\theta) R(\tau)]$, they suffer from high variance in gradient estimates, often resulting
 205 in slow convergence and suboptimal performance.

206 Prior in variance reduction techniques, such as SVRG and SPIDER, have addressed this challenge in
 207 supervised learning but require adaptation for non-stationary RL problems, where $p(\tau|\theta)$ changes
 208 with θ . Methods like HAPG and SRVR-PG use recursive gradient updates to achieve better sample
 209 complexity, reducing the number of trajectories required to find an ϵ -stationary point of $J(\theta)$.
 210 However, these methods often rely on large batch sizes and intricate learning rate schedules, limiting
 211 their practical efficiency.

212 To overcome these limitations, this paper introduces momentum-based policy gradient methods—IS-
 213 MBPG and HA-MBPG—which incorporate momentum-driven updates to stabilize gradient estima-
 214 tion and adapt learning rates dynamically. Momentum is a technique used in optimization to stabilize
 215 and accelerate convergence. In the context of policy gradient methods, it involves combining the
 216 current gradient estimate with a fraction of the previous update to reduce variance and smooth the
 217 optimization path. Formally, the momentum update can be represented as $u_t = \beta_t g_t + (1 - \beta_t) u_{t-1}$,
 218 where g_t is the current gradient, u_t is the momentum-adjusted gradient, and β_t controls the influence
 219 of the current and past gradients. This approach helps mitigate noise in stochastic updates and en-
 220 hances sample efficiency. These methods achieve the optimal sample complexity of $\mathcal{O}(\epsilon^{-3})$ without
 221 requiring large batches. Importantly, they maintain efficiency by updating parameters using only
 222 a single trajectory at each iteration, offering a significant advancement in reinforcement learning
 223 optimization.

224 4.2 IS-MBPG Algorithm Discussion and Theoretical Setup

225 The Important-Sampling Momentum-Based Policy Gradient algorithm combines momentum-based
 226 updates with importance sampling. At its core, the algorithm seeks to reduce the variance of
 227 stochastic gradients while maintaining computational efficiency. This is achieved by leveraging a
 228 momentum-driven variance reduction mechanism incorporating importance sampling to correct for
 229 non-stationarity in the underlying trajectory distribution $p(\tau|\theta)$.

230 The main optimization goal is to maximize the expected cumulative reward $J(\theta)$: $J(\theta) =$
 231 $\mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau)]$ where the gradient of $J(\theta)$ is approximated using the stochastic gradient:

232 $\hat{\nabla} J(\theta) = \frac{1}{|B|} \sum_{\tau \in B} g(\tau, \theta)$, and $g(\tau, \theta)$ is defined as: $g(\tau, \theta) = \sum_{h=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(a_h | s_h) \cdot$
 233 $\sum_{h=0}^{H-1} \gamma^h R(s_h, a_h)$, where π_{θ} is the policy, γ is the discount factor, and H is the horizon length.

234 The IS-MBPG algorithm refines the above gradient approximation by introducing a momentum-
 235 based update: $u_t = \beta_t g(\tau_t | \theta_t) + (1 - \beta_t) [u_{t-1} + g(\tau_t | \theta_t) - w(\tau_t | \theta_{t-1}, \theta_t) g(\tau_t | \theta_{t-1})]$, where
 236 $w(\tau_t | \theta_{t-1}, \theta_t)$ is the importance sampling weight: $w(\tau_t | \theta_{t-1}, \theta_t) = \frac{p(\tau_t | \theta_{t-1})}{p(\tau_t | \theta_t)} = \prod_{h=0}^{H-1} \frac{\pi_{\theta_{t-1}}(a_h | s_h)}{\pi_{\theta_t}(a_h | s_h)}$
 237 The weight $w(\tau_t | \theta_{t-1}, \theta_t)$ adjusts for the change in policy parameters between successive iterations,
 238 ensuring unbiased gradient estimates.

239 One of the key advantages of IS-MBPG is its ability to dynamically adapt learning rates based
 240 on the accumulated gradients. The learning rate η_t is computed as: $\eta_t = \frac{k}{(m + \sum_{i=1}^t G_i^2)^{1/3}}$, where
 241 $G_t = \|g(\tau_t, \theta_t)\|$ captures the magnitude of the gradient at each iteration. This adaptive learning
 242 rate ensures that the algorithm remains stable and converges efficiently, particularly in high-variance
 243 settings.

244 The momentum-based update in IS-MBPG reduces the variance of the gradient estimator through a
 245 combination of the several SGD variant techniques. When $\beta_t = 0$, the algorithm focuses on variance
 246 reduction. Conversely, when $\beta_t = 1$, the algorithm performs like SGD, prioritizing simplicity and
 247 speed. This flexibility allows IS-MBPG to balance exploration and exploitation dynamically. The
 248 IS-MBPG algorithm achieves the optimal sample complexity of $\mathcal{O}(\epsilon^{-3})$ for finding an ϵ -stationary
 249 point of the performance function $J(\theta)$. This improvement is achieved without relying on large
 250 batch sizes or double-loop structures, making IS-MBPG computationally efficient and scalable.
 251 Additionally, the algorithm only requires a single trajectory per iteration, significantly reducing the
 252 overhead compared to other variance-reduced policy gradient methods.

253 We present the convergence properties of the IS-MBPG algorithm. The theoretical guarantees of
 254 IS-MBPG rely on several key assumptions and propositions that establish its sample complexity and
 255 performance bounds. First make the following assumptions: The gradient and Hessian matrix of
 256 the function $\log \pi_{\theta}(a | s)$ are assumed to be bounded. That is, there exist constants M_g and $M_h > 0$
 257 such that: $\|\nabla_{\theta} \log \pi_{\theta}(a | s)\| \leq M_g$, $\|\nabla_{\theta}^2 \log \pi_{\theta}(a | s)\| \leq M_h$. Secondly the variance of the
 258 stochastic gradient $g(\tau | \theta)$ is bounded. Specifically, there exists a constant $\sigma > 0$ such that for all
 259 θ : $\mathbb{V}(g(\tau | \theta)) = \mathbb{E}\|g(\tau | \theta) - \nabla J(\theta)\|^2 \leq \sigma^2$. Finally, the variance of the importance sampling
 260 weight $w(\tau | \theta_1, \theta_2) = p(\tau | \theta_1) / p(\tau | \theta_2)$ is bounded. That is there exists a constant $W > 0$ such
 261 that: $\mathbb{V}(w(\tau | \theta_1, \theta_2)) \leq W$, $\forall \theta_1, \theta_2 \in \mathbb{R}^d$, $\tau \sim p(\tau | \theta_2)$. Note these assumptions are standard in
 262 the convergence analysis of policy gradient methods as seen in the previous paper with similar
 263 assumptions and ensure that the stochastic gradient and importance sampling weights remain stable.

264 We now present a few propositions. The first is Lipschitz Properties and Smoothness. Suppose $g(\tau | \theta)$
 265 is the policy gradient estimator. Under the boundedness assumptions, the following properties hold:
 266 The stochastic gradient $g(\tau | \theta)$ is L -Lipschitz differentiable, i.e., $\|g(\tau | \theta) - g(\tau | \theta')\| \leq L\|\theta - \theta'\|$,
 267 where $L = \frac{M_h R}{(1-\gamma)^2}$. Next, the performance function $J(\theta)$ is L -smooth, i.e., $\|\nabla^2 J(\theta)\| \leq L$. And
 268 lastly the stochastic gradient $g(\tau | \theta)$ is bounded, i.e., $\|g(\tau | \theta)\| \leq G$ for all $\theta \in \mathbb{R}^d$, where $G = \frac{M_g R}{(1-\gamma)^2}$.

269 The next proposition is a bounded Hessian: Under the same assumptions, for all θ , the Hessian
 270 estimator satisfies: $\|\nabla^2(\theta, \tau)\|^2 \leq \frac{H^2 M_g^4 R^2 + M_h^2 R^2}{(1-\gamma)^4} = \tilde{L}^2$. This implies that $J(\theta)$ is \tilde{L} -smooth.

271 4.3 Convergence Analysis of IS-MBPG

272 **Theorem 4: Convergence of IS-MBPG.** Let $\{\theta_t\}_{t=1}^T$ be the sequence generated by the IS-MBPG
 273 algorithm. Set $k = \mathcal{O}(G^{2/3}/L)$, $c = G^2/(3k^3L) + 104B^2$, $m = \max\{2G^2, (2Lk)^3, (ck^2/L)^3\}$,
 274 and $\eta_0 = k/m^{1/3}$. Then, we have:

$$\mathbb{E}\|\nabla J(\theta_{\zeta})\| \leq \frac{\sqrt{2\Omega m^{1/6}}}{\sqrt{T}} + \frac{2\Omega^{3/4}}{\sqrt{T}} + \frac{2\sqrt{\Omega}\sigma^{1/3}}{T^{1/3}},$$

275 where $\Omega = \frac{1}{k} \left(16(J^* - J(\theta_1)) + \frac{m^{1/3}\sigma^2}{8B^2k} + \frac{c^2k^3}{4B^2} \ln(T+2) \right)$ and $J^* = \sup_{\theta} J(\theta) < \infty$.

276 We now present the proof sketch for this. First define two lemmas which will be used in the proof:

277 **Lemma 3.** Under Assumption 1, let $e_t = \nabla J(\theta_t) - u_t$. Given: $0 < \eta_t \leq \frac{1}{2L}$ for all $t \geq 1$ we have

$$E[J(\theta_{t+1})] \geq E \left[J(\theta_t) - \frac{3\eta_t}{4} \|e_t\|^2 + \frac{\eta_t}{8} \|\nabla J(\theta_t)\|^2 \right].$$

278 **Lemma 4. Variance Dynamics of the Stochastic Gradient.** Assume that the stochastic policy
279 gradient u_t is generated by Algorithm 1, and let $e_t = u_t - \nabla J(\theta_t)$. Then the following inequality
280 holds:

$$\mathbb{E} [\eta_{t-1}^{-1} \|e_t\|^2] \leq 2\beta_t^2 \eta_{t-1}^{-1} G_t^2 + \eta_{t-1}^{-1} (1 - \beta_t)^2 (1 + 8\eta_{t-1}^2 B^2) \mathbb{E} \|e_{t-1}\|^2 + 8(1 - \beta_t)^2 B^2 \eta_{t-1} \|\nabla J(\theta_{t-1})\|^2,$$

281 where $B^2 = L^2 + 2G^2 C_w^2$ and $C_w = \sqrt{H(2HM_g^2 + M_h)(W + 1)}$. The proof for this ex-
282 pands $\mathbb{E}[\eta_{t-1}^{-1} \|e_t\|^2]$. The proof of this lemma comes down to bounding $\mathbb{E}[\eta_{t-1}^{-1} \|k_{etk}\|^2] \leq 2\eta_{t-1}^{-1} (1 -$
283 $\beta_t)^2 \mathbb{E} \|k_{et-1}\|^2 + 2\beta_t^2 \eta_{t-1}^{-1} G_t^2 + 2(1 - \beta_t)^2 \eta_{t-1}^{-1} \mathbb{E} \|g(\tau_t|\theta_t) - w(\tau_t|\theta_{t-1}, \theta_t)g(\tau_t|\theta_{t-1})\|^2$. Denote
284 the last expectation as T_1 . Further bounding on T_1 yields $T_1 \leq 2(L^2 + 2G^2 C_w^2) \|\theta_t - \theta_{t-1}\|^2$.

285 We now cover the main convergence proof of Theorem 4 in the paper. Assume the given information
286 in the start of the theorem. We first do parameter initialization and provide bounds on η_t and β_t .

287 We begin by noting that $m \geq (2Lk)^3$ ensures $\eta_t \leq k/m^{1/3} \leq 1/2L$. Similarly, $m \geq (ck/2L)^3$
288 ensures $\beta_{t+1} = c\eta_t^2 \leq c\eta_t/2L \leq \frac{ck}{2Lm^{1/3}} \leq 1$. Thus, the learning rate η_t and momentum parameter
289 β_t remain valid for all t .

290 From Lemma 4, the error term $e_t = u_t - \nabla J(\theta_t)$ satisfies: $\mathbb{E} [\eta_{t-1}^{-1} \|e_t\|^2 - \eta_{t-2}^{-1} \|e_{t-1}\|^2] \leq$
291 $\mathbb{E} [2\beta_t^2 \eta_{t-1}^{-1} G_t^2 + \eta_{t-1}^{-1} (1 - \beta_t)^2 (1 + 8\eta_{t-1}^2 B^2) \|e_{t-1}\|^2 + 8(1 - \beta_t)^2 B^2 \eta_{t-1} \|\nabla J(\theta_{t-1})\|^2]$.

292 After we proceed with refining terms and bounding $\eta_{t-1}^{-1} - \eta_{t-2}^{-1}$. The term $\eta_{t-1}^{-1} - \eta_{t-2}^{-1}$ is
293 bounded by exploiting the concavity of $x^{1/3}$. Specifically: $\eta_{t-1}^{-1} - \eta_{t-2}^{-1} \leq \frac{G_t^2}{3k(m + \sum_{i=1}^{t-1} G_i^2)^{2/3}} \leq$
294 $\frac{G_t^2}{3k(m + \sum_{i=1}^t G_i^2)^{2/3}}$. Using $m \geq 2G^2$, this implies: $\eta_{t-1}^{-1} - \eta_{t-2}^{-1} \leq \frac{G_t^2}{3k^3 L} \eta_t$.

295 Next we bound on the auxiliary terms. Substituting the previous result into the upper bound of the error
296 dynamics: $T_2 = (\eta_{t-1}^{-1} - \eta_{t-2}^{-1} + 8B^2 \eta_t - \beta_t \eta_{t-1}^{-1} - 8\eta_t \beta_t B^2) \|e_t\|^2$. Using $c = \frac{G^2}{3k^3 L} + 104B^2$,
297 we find: $T_2 \leq -96B^2 \eta_t \|e_t\|^2$.

298 Now define a Lyapunov function, a scalar function that is used in control theory to prove the stability
299 of an equilibrium of a first order ode. $\Phi_t = J(\theta_t) - \frac{1}{128B^2 \eta_{t-1}} \|e_t\|^2$. Using Lemma 3, the analysis

300 yields: $\mathbb{E}[\Phi_{t+1} - \Phi_t] \geq \mathbb{E} \left[-\frac{c\eta_t^3 G_{t+1}^2}{64B^2} + \frac{\eta_t}{16} \|\nabla J(\theta_t)\|^2 \right]$. Summing over $t = 1$ to T , this implies:
301 $\sum_{t=1}^T \mathbb{E}[\eta_t \|\nabla J(\theta_t)\|^2] \leq \mathbb{E} \left[16(J^* - J(\theta_1)) + \frac{1}{8B^2 \eta_0} \|e_1\|^2 + \frac{c^2 k^3}{4B^2} \ln(T + 2) \right]$.

302 Applying the Cauchy-Schwarz inequality, we obtain: $\frac{1}{T} \sum_{t=1}^T \mathbb{E} \|\nabla J(\theta_t)\| \leq$
303 $\sqrt{\frac{1}{T} \sum_{t=1}^T \mathbb{E} \|\nabla J(\theta_t)\|^2}$. Combining with the above, and using Ω as defined, we conclude:
304 $\frac{1}{T} \sum_{t=1}^T \mathbb{E} \|\nabla J(\theta_t)\| \leq \sqrt{2\Omega m^{1/6}} + \frac{2\Omega^{3/4}}{\sqrt{T}} + \frac{2\sqrt{\Omega} \sigma^{1/3}}{T^{1/3}}$. This establishes the theorem.

305 **Remark 1: Sample Complexity.** Since $\Omega = \mathcal{O}(\ln T)$, Theorem 1 shows that the IS-MBPG
306 algorithm achieves an $\mathcal{O}(\sqrt{\ln T}/T^{1/3})$ convergence rate. The algorithm requires a single trajectory
307 per iteration, and the total sample complexity is $\mathcal{O}(\epsilon^{-3})$ to achieve an ϵ -stationary point, where
308 $T = \epsilon^{-3}$.

309 4.4 HA-MBPG Algorithm Discussion and Analysis

310 The Hessian-Aided Momentum-Based Policy Gradient (HA-MBPG) algorithm builds upon the
311 momentum-based updates of IS-MBPG while introducing a Hessian-aided technique to further refine
312 the gradient estimates with the same theoretical assumptions as before. The core idea is to leverage
313 second-order information, approximated through the Hessian of the policy's performance function
314 $J(\theta)$, to improve convergence speed and reduce variance in the stochastic gradient estimates.

The optimization objective remains the maximization of the expected cumulative reward: $J(\theta) = \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau)]$. The HA-MBPG algorithm incorporates the Taylor expansion to approximate the difference in gradients $\nabla J(\theta_t)$ and $\nabla J(\theta_{t-1})$ as: $\nabla J(\theta_t) - \nabla J(\theta_{t-1}) \approx \left[\int_0^1 \nabla^2 J(\theta_{t-1} + \alpha(\theta_t - \theta_{t-1})) d\alpha \right] (\theta_t - \theta_{t-1})$, where $\nabla^2 J(\theta)$ represents the Hessian matrix of $J(\theta)$. This term captures the curvature of the performance function and provides additional information to guide the update process.

To operationalize this, the algorithm constructs an unbiased estimate of the gradient difference using the sampled trajectory τ and an auxiliary random variable $\alpha \sim U[0, 1]$: $\Delta_t = \nabla^2 J(\theta_{t-1} + \alpha(\theta_t - \theta_{t-1}))(\theta_t - \theta_{t-1})$. In practice, the exact computation of $\nabla^2 J(\theta)$ is computationally expensive and thus intractable. Instead, HA-MBPG approximates the Hessian-vector product $\nabla^2 J(\theta)v$ using a finite-difference method from numerical PDE solving methods: $\nabla^2 J(\theta)v \approx \frac{\nabla J(\theta + \delta v) - \nabla J(\theta - \delta v)}{2\delta}$, where $\delta > 0$ is a small scalar and $v = \theta_t - \theta_{t-1}$. This approximation balances computational efficiency with accuracy, making the Hessian-aided updates feasible for practical implementation.

The gradient update rule in HA-MBPG incorporates both momentum and the Hessian-aided term: $u_t = \beta_t w(\tau|\theta_t, \theta_{t-1})g(\tau|\theta_t) + (1 - \beta_t)[u_{t-1} + \Delta_t]$, where $w(\tau|\theta_t, \theta_{t-1})$ is the importance sampling weight that adjusts for changes in policy parameters: $w(\tau|\theta_t, \theta_{t-1}) = \frac{p(\tau|\theta_t)}{p(\tau|\theta_{t-1})}$.

The learning rate η_t is adaptively calculated as: $\eta_t = \frac{k}{(m + \sum_{i=1}^t G_i^2)^{1/3}}$, where $G_t = \|g(\tau_t, \theta_t)\|$ reflects the magnitude of the gradient at iteration t . This adaptive rate ensures stability and efficient convergence by scaling learning rates based on the accumulated gradients. We now provide a analysis of the convergence proof.

4.5 Convergence Analysis of HA-MBPG

We provide the following convergence theorem and the lemma used in the proof. Note the proof similarity to the first order momentum based approach first bounding recurrence of the error term using the 2nd order analog to lemma 4, lemma 5, and then defining a Lyapunov potential function, to use smoothness properties to telescope between iterative potential functions and then using Cauchy Schwartz to bound the expected gradient norm.

Lemma 5. Assume that the stochastic policy gradient u_t is generated by Algorithm 2. Let $e_t = u_t - \nabla J(\theta_t)$, then the following holds:

$$\mathbb{E} [\eta_{t-1}^{-1} \|e_t\|^2] \leq 4(W+1)\beta_t^2 \eta_{t-1}^{-1} G_t^2 + \eta_{t-1}^{-1} (1 - \beta_t)^2 (1 + 4\eta_{t-1}^2 L^2) \mathbb{E} \|e_{t-1}\|^2 + 4(1 - \beta_t)^2 L^2 \eta_{t-1} \|\nabla J(\theta_{t-1})\|^2.$$

Theorem 2: Convergence of HA-MBPG Algorithm. Let $\{\theta_t\}_{t=1}^T$ be the sequence generated by Algorithm 2, and set $k = \mathcal{O}(G^2/3L)$, $c = \frac{G^2}{3k^3L} + 52L^2$, $m = \max\{2G^2, (2Lk)^3, (ck/2L)^3\}$, and $\eta_0 = k/m^{1/3}$. Then, the following bound holds:

$$\mathbb{E} \|\nabla J(\theta_\zeta)\| \leq \sqrt{2\Lambda m^{1/6}} + \frac{2\Lambda^{3/4}}{\sqrt{T}} + \frac{2\sqrt{\Lambda}\sigma^{1/3}}{T^{1/3}},$$

where

$$\Lambda = \frac{1}{k} \left(16(J^* - J(\theta_1)) + \frac{m^{1/3}\sigma^2}{4L^2k} + \frac{(W+1)c^2k^3}{2L^2} \ln(T+2) \right),$$

and $J^* = \sup_\theta J(\theta) < \infty$.

Starting off, given $m \geq (2Lk)^3$, we have $\eta_t \leq k/m^{1/3} \leq 1/2L$. Similarly, $m \geq (ck/2L)^3$ ensures $\beta_{t+1} = c\eta_t^2 \leq c\eta_t/2L \leq \frac{ck}{2Lm^{1/3}} \leq 1$. Thus, both the learning rate η_t and momentum coefficient β_t remain valid for all t . Using Lemma 5, the error term $e_t = u_t - \nabla J(\theta_t)$ satisfies the following recurrence: $\mathbb{E} [\eta_{t-1}^{-1} \|e_t\|^2 - \eta_{t-2}^{-1} \|e_{t-1}\|^2] \leq \mathbb{E} \left[4(W+1)\beta_t^2 \eta_{t-1}^{-1} G_t^2 + \eta_{t-1}^{-1} (1 - \beta_t)^2 (1 + 4\eta_{t-1}^2 L^2) \|e_{t-1}\|^2 + 4(1 - \beta_t)^2 L^2 \eta_{t-1} \|\nabla J(\theta_{t-1})\|^2 \right]$. Using the concavity of $x^{1/3}$, we bound the difference $\eta_{t-1}^{-1} - \eta_{t-2}^{-1}$: $\eta_{t-1}^{-1} - \eta_{t-2}^{-1} = \frac{1}{k} \left((m + \sum_{i=1}^t G_i^2)^{1/3} - (m + \sum_{i=1}^{t-1} G_i^2)^{1/3} \right) \leq \frac{G_t^2}{3k(m + \sum_{i=1}^t G_i^2)^{2/3}}$. Using $m \geq 2G^2$, we find: $\eta_{t-1}^{-1} - \eta_{t-2}^{-1} \leq \frac{G_t^2}{3k^3L} \eta_t$. Substituting the above

355 result into the error dynamics: $T_4 = (\eta_{t-1}^{-1} - \eta_{t-2}^{-1} + 4L^2\eta_t - \beta_t\eta_{t-1}^{-1} - 4\eta_t\beta_tL^2) \|e_t\|^2$. Using
356 $c = \frac{G^2}{3k^3L} + 52L^2$, we find: $T_4 \leq -48L^2\eta_t\|e_t\|^2$.
357 Define the Lyapunov function $\Psi_t = J(\theta_t) - \frac{1}{64L^2\eta_{t-1}}\|e_t\|^2$. Using the smoothness of $J(\theta)$ and the re-
358 fined recurrence, we derive: $\mathbb{E}[\Psi_{t+1} - \Psi_t] \geq \mathbb{E}\left[-\frac{(W+1)c^2\eta_t^3G_{t+1}^2}{32L^2} + \frac{\eta_t}{16}\|\nabla J(\theta_t)\|^2\right]$. Summing over
359 $t=1$ to T , $\sum_{t=1}^T \mathbb{E}[\eta_t\|\nabla J(\theta_t)\|^2] \leq \mathbb{E}\left[16(J^* - J(\theta_1)) + \frac{1}{4L^2\eta_0}\|e_1\|^2 + \frac{(W+1)c^2k^3}{2L^2}\ln(T+2)\right]$.
360 Using the Cauchy-Schwarz inequality, the squared gradient norm is bounded by:
361 $\frac{1}{T}\sum_{t=1}^T \mathbb{E}\|\nabla J(\theta_t)\| \leq \sqrt{\frac{1}{T}\sum_{t=1}^T \mathbb{E}\|\nabla J(\theta_t)\|^2}$. Substituting the above and simplifying us-
362 ing Λ : $\mathbb{E}\|\nabla J(\theta_t)\| \leq \sqrt{2\Lambda m^{1/6}} + \frac{2\Lambda^{3/4}}{\sqrt{T}} + \frac{2\sqrt{\Lambda}\sigma^{1/3}}{T^{1/3}}$. Since $\Lambda = \mathcal{O}(\ln T)$, the convergence rate is
363 $\mathcal{O}(\sqrt{\ln T}/T^{1/3})$. Setting $T = \epsilon^{-3}$, the sample complexity is $\mathcal{O}(\epsilon^{-3})$, completing the proof.

364 5 Algorithm Comparison

365 The main similarity between these papers is that they all discuss methods for finding ϵ -stationary
366 points of a policy π_θ 's objective function $J(\theta)$. These papers both describe their algorithm, as well as
367 provide proofs of their algorithm's convergence guarantees.

368 In the SRVR-PG and RPG papers, the optimization relies on first-order methods, which utilize
369 estimates of the policy gradient $\nabla J(\theta)$ without explicitly incorporating higher-order derivatives
370 such as the Hessian $\nabla^2 J(\theta)$. For instance, the SRVR-PG method introduces a variance reduction
371 technique by maintaining a control variate, which stabilizes gradient estimates and achieves a sample
372 complexity of $\mathcal{O}(\epsilon^{-3})$. This is done by iteratively estimating gradients through stochastic recursive
373 updates, a process governed by: $\nabla_{\text{SVR}} J(\theta) = \nabla J(\theta_k) + \frac{1}{B} \sum_{i=1}^B (\nabla J_i(\theta) - \nabla J_i(\theta_k))$, where B is
374 the mini-batch size, and $\nabla J_i(\cdot)$ denotes the stochastic gradient of the i -th trajectory. This method
375 avoids computing second-order derivatives, making it computationally attractive.

376 The RPG paper extends the focus to robust reinforcement learning under model mismatch. Here,
377 the objective is to optimize the worst-case performance of a policy over an uncertainty set \mathcal{P} . The
378 robust policy gradient is derived using a sub-gradient approach due to the non-differentiability of the
379 worst-case value function: $\psi_\rho(\theta) = \sum_{s \in \mathcal{S}} d_\rho^{\pi_\theta}(s) \sum_{a \in \mathcal{A}} \nabla \pi_\theta(a|s) Q^{\pi_\theta}(s, a)$, where $d_\rho^{\pi_\theta}(s)$ is the
380 discounted visitation distribution and $Q^{\pi_\theta}(s, a)$ incorporates the worst-case transition kernel. The
381 RPG algorithm employs smoothing techniques to approximate the max operator, enabling efficient
382 computation of gradients while maintaining theoretical robustness guarantees.

383 On the other hand, the MBPG paper incorporates second-order information by explicitly using the
384 Hessian $\nabla^2 J(\theta)$ in its momentum-based updates. This additional information allows the algorithm to
385 adaptively adjust the update direction, leading to accelerated convergence. The method constructs a
386 momentum term that blends the gradient and Hessian information: $M_t = \beta M_{t-1} + \nabla^2 J(\theta_t) \nabla J(\theta_t)$,
387 where β is the momentum parameter. This approach achieves superior theoretical convergence rates,
388 particularly in non-convex settings, by leveraging curvature information to navigate the landscape
389 of $J(\theta)$ more effectively. The paper proves that MBPG achieves a sample complexity of $\mathcal{O}(\epsilon^{-3})$ in
390 certain settings, outperforming purely first-order methods.

391 While all three papers share a common goal of optimizing $J(\theta)$ efficiently, their methodological
392 choices reflect a trade-off between computational simplicity and theoretical guarantees. SRVR-PG
393 and RPG prioritize sample efficiency and robustness using first-order gradients, while MBPG achieves
394 faster convergence by incorporating second-order derivatives at the cost of increased computational
395 complexity.

396 6 Conclusion

397 In this report, we describe three fundamental papers for efficient policy gradient algorithms for
398 convergence to ϵ -stationary policies. We also present the convergence analyses for all of these
399 methods, and compare and contrast these algorithms in terms of their efficiency in converging to an
400 ϵ -stationary policy. We hope that our work provides inspiration to future analyses of policy gradient
401 algorithms.

402 **References**

- 403 [1] Xu, Pan & Gao, Felicia & Gu, Quanquan (2020) Sample Efficient Policy Gradient Methods With Recursive
404 Variance Reduction, *International Conference on Learning Representations*
- 405 [2] Wang, Yue & Zou, Shaofeng (2022) Policy Gradient Method For Robust Reinforcement Learning, *International Conference on Machine Learning*
406
- 407 [3] Huang, Feihu & Gao, Shangqian & Pei, Hian & Huang, Heng (2020) Momentum-Based Policy Gradient
408 Methods, *International Conference on Machine Learning*

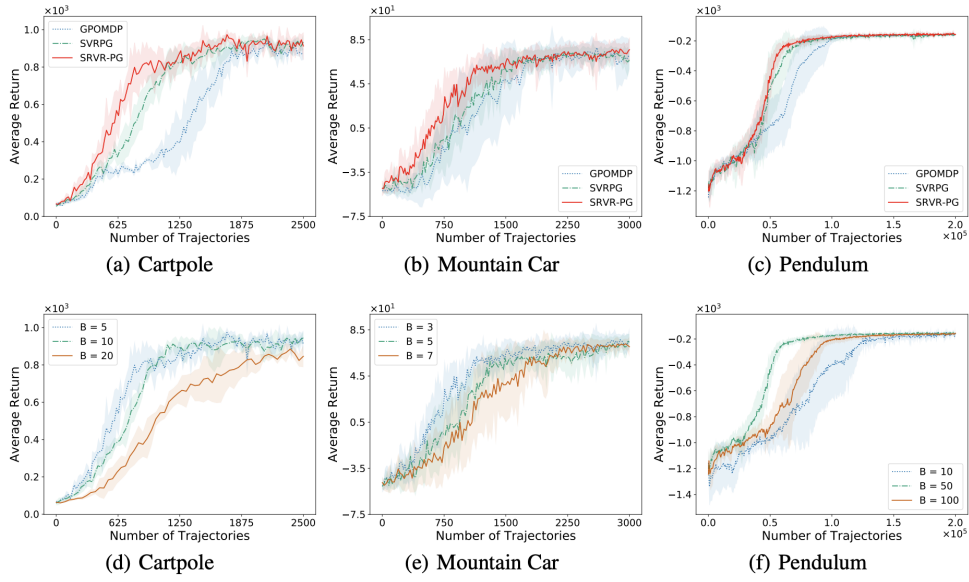


Figure 1: (a)-(c): Comparison of different algorithms. Experimental results are averaged over 10 repetitions. (d)-(f): Comparison of different batch size B on the performance of SRVR-PG.

Figure 1: SRVR-PG Experiments Results.

A Appendix / supplemental material

Shown above are the experimental results for the SRVR-PG paper.